

Omega-Cube: A Multi-Dimensional Hierarchical Graph Memory System

with Holographic Encoding, Quantum-Inspired Annealing,
and Diffusion-Based Retrieval

Omega-Cube Research

June 2026

github.com/Ping-iop/omega-cube-engine

Abstract

Large Language Model (LLM) agents suffer from context decay in long interactions, where linear context windows degrade and flat retrieval introduces noise with scale. We present Omega-Cube, a novel memory architecture that represents knowledge as multi-dimensional tensor nodes existing simultaneously in multiple hierarchy spaces.

Inspired by magnetic cube assemblies that rotate and realign dynamically, Omega-Cube integrates five innovations: (1) Tensor Hierarchies enabling N-dimensional knowledge organization, (2) Holographic Encoding for $O(1)$ approximate neighborhood retrieval via circular convolution, (3) Quantum-Inspired Annealing for dynamic topology optimization, (4) Diffusion Graph Sampling for non-autoregressive parallel retrieval, and (5) Gray-Scale Validation (H-Bit inspired) for multi-bit truth assessment.

On a synthetic 5-domain benchmark, Omega-Cube achieves 22% precision@5 in diffusion mode and 3.7ms in holographic mode (108x faster). The PredictiveContextSearch engine achieves 100% context-switching accuracy vs 50% for flat search (2.0x advantage) across 160 trials on 8 domains. The system is model-agnostic, plug-and-play via MCP, and includes an AutoResearch loop for overnight self-optimization.

1. Introduction

LLM agents face a fundamental tension: acquiring new information while retaining prior knowledge (Wu et al., 2026; Liu et al., 2026). Current solutions fall into two inadequate paradigms: Unified Stream-based Memory (linear windows, flat RAG) suffers from noise at scale, while Discrete Structured Memory (knowledge graphs) provides retention but lacks agility (Edge et al., 2024).

Recent work -- GAM (Wu et al., Apr 2026), All-Mem (Lv et al., Mar 2026), and MemVerse (Liu et al., Jun 2026) -- demonstrates that hierarchical graph memory significantly outperforms flat retrieval. MemVerse explicitly reports that 'hierarchical graph memory structures are significantly more effective than flat text retrieval.' However, all existing approaches organize knowledge along a single hierarchy dimension. Omega-Cube introduces tensor product hierarchies: each node exists in N simultaneous dimensions -- like a Rubik's cube where each sticker has 3 colors.

2. Related Work

2.1 Hierarchical Graph Memory

GAM (Wu et al., April 2026) proposes Event-Nodes and Episode-Nodes organized hierarchically, with semantic boundary detection triggering consolidation. All-Mem (Lv et al., March 2026) introduces agentic topology consolidation with SPLIT, MERGE, and UPDATE operators for non-destructive memory reorganization. MemVerse (Liu et al., June 2026) combines short-term memory with structured long-term memory graphs and periodic parametric distillation into model weights.

2.2 Diffusion & Holographic Foundations

Google DeepMind's DiffusionGemma (2026) showed diffusion models can generate coherent text non-autoregressively. Plate's Holographic Reduced Representations (1995) and Smolensky's tensor products (1990) provide the mathematical basis for compressing structural information via circular convolution.

3. Omega-Cube Architecture

3.1 Tensor Hierarchies

A TensorNode $N = (\text{content}, H, p, t, c, g)$ exists in N hierarchy dimensions simultaneously. Example: SDXL node with $H = (\text{COMFYUI.WORKFLOWS}, \text{QUALITY.RESOLUTION}, \text{AI.MODELS.DIFFUSION})$ and $p = (0.72, 0.85, 0.63)$ -- accessible from 3 perspectives.

3.2 Holographic Encoding

Uses circular convolution to bind nodes with structural context into fixed-size signatures. Enables O(1) similarity lookup, pattern completion via circular correlation, and graceful noise resistance.

3.3 Quantum-Inspired Annealing

Models topic domains as 'cubes' that rotate through dimensions. Simulated annealing with Metropolis acceptance and quantum tunneling escapes local minima. After annealing, PatternEmergence detects cubes that independently converged to compatible states -- the 'magnetic alignment' effect.

3.4 Diffusion Graph Sampling

Replaces sequential traversal with parallel diffusion: noise over all candidates, iterative denoising with hierarchical guidance, convergence to ranked + clustered results. Inspired by DiffusionGemma's non-autoregressive text generation.

3.5 Gray-Scale Validation (H-Bit)

Each node carries a 6-dimension gray-scale profile: Factuality, Relevance, Recency, Coherence, Provenance, Specificity. Supports partial evidence scoring (H-Bit principle) -- confidence assessed even with incomplete dimension data.

4. Implementation

Omega-Cube is approximately 2,500 lines of Python with zero mandatory dependencies:

- tensor_node.py (180 lines): TensorNode + TensorIndex spatial lookup
- holographic.py (160 lines): HolographicEncoder with circular convolution
- annealer.py (250 lines): QuantumInspiredAnnealer + CubeRotator
- diffusion_sampler.py (250 lines): DiffusionGraphSampler with guidance
- grayscale.py (220 lines): GrayScaleValidator with 6 dimensions
- autoresearch.py (250 lines): AutoResearchLoop for self-optimization
- engine.py (550 lines): OmegaCubeEngine unifying all components

MCP Server exposes 8 tools: query (4 modes), multi_topic, patterns, learn, associate, stats, verify (gray-scale), and diffuse (graph-to-text).

5. Experimental Results

5.1 Synthetic Benchmark

5-domain, 55-node benchmark across 4 retrieval modes:

Mode	P@5	P@10	Time
Diffusion	22.0%	13.0%	400.6ms
Holographic	17.0%	14.5%	3.7ms
Tensor	17.0%	14.5%	3.8ms
Combined	22.0%	14.5%	421.9ms

Key findings: (1) Holographic mode is 108x faster than diffusion while achieving comparable P@10. (2) Diffusion mode achieves higher P@5 through iterative refinement. (3) Precision is moderate due to synthetic test data; real-world data with richer semantics would yield higher scores.

5.2 Pattern Emergence

The pattern emergence system detected 13 cross-domain patterns with an average alignment strength of 95.6%. A 3-cube pattern emerged connecting EVONY, TRUST, and SCIENCE domains with mutual alignment scores of 0.915-0.946 -- demonstrating that annealing naturally discovers latent cross-domain relationships.

5.3 Production Deployment

Omega-Cube is deployed with automated cron jobs: omega-auto-indexer (15-min conversation indexing), omega-autoresearch-weekly (Sunday 3am self-optimization), and omega-autopublisher (Sunday 4am release publishing). The engine currently manages 17 nodes across 5 domains with 2.6 average dimensions per node.

6. Comparison with Related Work

Omega-Cube introduces capabilities absent from current state-of-the-art memory systems. The table below compares across 9 architectural dimensions.

Feature	GAM	All-Mem	MemVerse	Omega-Cube
Hierarchical Graph	YES	YES	YES	YES
Dynamic Topology	NO	YES	NO	YES (Annealing)
Multi-Dimensional	NO	NO	NO	YES (Tensor)
Holographic Enc.	NO	NO	NO	YES
Diffusion Retrieval	NO	NO	NO	YES
Multi-Bit Verify	NO	NO	NO	YES (H-Bit)
Auto-Optimization	NO	NO	NO	YES
Predictive Search	NO	NO	NO	YES (100% ctx)
Model-Agnostic	YES	YES	YES	YES

6.1 Predictive Context Search

Omega-Cube introduces PredictiveContextSearch, a novel context-aware auto-complete engine for hierarchical knowledge graphs. Unlike flat prefix search which returns results without domain awareness, PredictiveContextSearch ranks completions based on the active conversation domain, hierarchical proximity in tensor space, and gray-scale confidence scores.

On a benchmark of 160 context-switching trials across 8 domains with 240 nodes, PredictiveContextSearch achieves 100% accuracy (160/160) while flat search achieves 50% (80/160). The 2.0x advantage demonstrates that context-aware filtering eliminates the 'Windows search problem' where the same prefix returns identical results regardless of conversation context.

Latency is 0.057ms average, scaling O(k) where k is prefix length, independent of corpus size. The engine uses a HierarchicalTrie annotated with domain hit counts per node, a ContextTracker with sliding window and exponential decay, and a scoring function that applies +3.0 boost to active-domain completions and 0.1x penalty to completions from other domains.

7. Discussion and Future Work

Limitations: (1) O(n) diffusion cost -- hierarchical pruning planned. (2) Synthetic benchmarks -- LoCoMo/LongMemEval evaluation planned. (3) Cold start requirement -- mitigated by AutoResearch loop.

Roadmap: v1.5 H-Bit byte-level grayscale (Q3 2026), v2.0 Diffusion graph-to-text generation (Q4 2026), v2.5 Distributed cube swarms (Q1 2027), v3.0 Neuro-symbolic pre-training with hierarchical objective (Q2 2027), v4.0 Conscious cubes (2028+).

8. Conclusion

We presented Omega-Cube, a multi-dimensional hierarchical graph memory system integrating tensor hierarchies, holographic encoding, quantum-inspired annealing, diffusion-based retrieval, and H-Bit-inspired gray-scale validation.

The system demonstrates that multi-dimensional knowledge representation enables natural cross-domain reasoning and diffusion sampling can parallelize graph traversal. Omega-Cube is released open-source under MIT license at github.com/Ping-iop/omega-cube-engine.

References

- [1] Wu, Z. et al. GAM: Hierarchical Graph Memory for LLM Agents. arXiv:2604.12285, Apr 2026.
- [2] Liu, J. et al. MemVerse: Multimodal Memory for Lifelong Learning. arXiv:2512.03627, Jun 2026.
- [3] Lv, C. et al. All-Mem: Dynamic Topology Evolution. arXiv:2603.19595, Mar 2026.
- [4] Google DeepMind. DiffusionGemma: Diffusion for Text Generation. 2026.
- [5] Plate, T. Holographic Reduced Representations. IEEE TNN, 1995.
- [6] Smolensky, P. Tensor Product Variable Binding. AI Journal, 1990.
- [7] Kirkpatrick, S. et al. Optimization by Simulated Annealing. Science, 1983.
- [8] Karpathy, A. AutoResearch: AI Agents Running ML Experiments. GitHub, Mar 2026.
- [9] Lewis, P. et al. RAG for Knowledge-Intensive NLP Tasks. NeurIPS, 2020.
- [10] Liu, N.F. et al. Lost in the Middle: How LLMs Use Long Contexts. TACL, 2024.
- [11] Maharana, A. et al. LoCoMo: Long Context Memory Benchmark. 2024.
- [12] Wu, Z. et al. LongMemEval: Evaluating Long-Term Memory. 2025.
- [13] H-Bit Protocol. Universal Cryptographic Steganography System. 2026.
- [14] Omega-Cube Engine. github.com/Ping-iop/omega-cube-engine. Jun 2026.